



# Form Service Schema - Business and Product Capability Analysis

## Supported Product Capabilities

### Reusability Model

The schema strongly emphasizes reuse at multiple levels. It defines **reusable field definitions** (`field_def`) that are tenant-scoped and can be used across different forms or even entities like tickets <sup>1</sup>. Each field definition has a stable key (`field_key`) unique within a tenant for consistent reference by APIs or forms <sup>2</sup>. This allows the same field (e.g. a "Priority" or "Customer Email" field) to be defined once and reused in many contexts without duplication. Field definitions are stored with both a *business key* and a *version*, meaning the concept of a field can evolve over time while retaining a consistent identity <sup>3</sup>.

Beyond individual fields, the schema supports **reusable components** – logical groupings of fields that can be treated as a single form element. The `component` table represents such a reusable component, which might be a composite widget or a sub-form (for example, an "Address block" containing multiple fields). Components are also tenant-scoped catalog entries and can originate from different sources (provider-supplied defaults, third-party marketplace, or tenant-defined) <sup>4</sup>. Like fields, components have stable keys and versioning (using a `component_business_key` and version) to manage iterations <sup>5</sup> <sup>6</sup>. This model enables a high degree of reuse: a tenant can maintain a library of common fields and components and assemble them into forms as needed, rather than building every form from scratch.

Notably, the schema treats forms themselves as definitional artifacts, separate from their usage. A form definition (`form` table) can also have a business key and version, implying that a form is a versioned entity (e.g. you could have "Onboarding Form v1", "v2", etc.) <sup>7</sup> <sup>8</sup>. Old versions can be retained (for compatibility or record-keeping) while new versions are published, which is a sophisticated capability not all CRM form builders expose.

### Composition Patterns

The form schema allows complex composition of form elements through a panel and nesting structure. An **actual form** is composed of one or more panels (`form_panel` entries), which act as containers or sections within the form <sup>9</sup>. Each panel can directly contain fields (`form_panel_field`) and can also contain component instances (`form_panel_component`) that embed a reusable component into the form <sup>10</sup> <sup>11</sup>. This design supports rich form layouts and hierarchies: for example, a form panel could have a mix of individual fields and a composite component (with its own internal panels and fields) placed in it.

Components themselves have an internal structure of panels (`component_panel`) which can nest hierarchically <sup>12</sup>. Within a component's panels, fields are placed via `component_panel_field` records. The schema even allows panels to be nested inside other panels within a component (via a parent panel reference), enabling multi-level grouping (like sub-sections inside a reusable component) <sup>13</sup>. However, the

schema does **not** indicate that components can directly include other components (there is no `component_panel_component` table), so nesting appears limited to fields and panels within a component. This is likely a deliberate design choice to keep the hierarchy manageable (forms can contain components, which contain fields, but you wouldn't embed a component inside another component directly in the catalog).

The composition approach is very flexible. It allows, for instance, building a complex form by assembling smaller pieces (fields and components) rather than designing monolithically. Each component placed on a form is given an `instance_key` that uniquely identifies that instance within the form <sup>14</sup>. This means even if the same component is used twice in one form, or the same field definition appears multiple times, each placement has a distinct key for tracking. Indeed, the schema explicitly accounts for **the same field definition appearing multiple times** in a form or across nested components – a scenario handled by generating a **fully qualified field path** for each field instance in submissions <sup>15</sup>. This path includes the form's key, panel keys, component instance keys, and field keys, ensuring every data entry can be traced back to its position in the form structure <sup>16</sup>. The flat submission storage model (one row per field instance) uses this path to avoid ambiguity between repeated fields <sup>15</sup>. This level of instance identification shows a robust composition pattern where forms are trees of panels -> fields/components, and every node in that tree is referenceable.

The schema also supports **override and configuration at each composition level**. Both field placements and component placements allow JSON-based `ui_config` or override data. For example, when a field is placed on a component panel, there is a `component_panel_field.ui_config` for overrides specific to that placement <sup>17</sup> <sup>18</sup>. Similarly, when a component is embedded in a form, the `form_panel_component` has an `ui_config` for that instance, and even a `nested_overrides` JSON that can override settings deep inside the component for that particular usage <sup>19</sup> <sup>20</sup>. This capability implies forms can be finely tuned – e.g. a component might normally have a field labeled "Name", but in one specific form you could override that label to "Full Name" without affecting the component globally. The presence of a **deterministic JSON merge utility** in the schema <sup>21</sup> <sup>22</sup> reinforces that overriding and merging configurations is a first-class concept. This level of compositional flexibility is a strength for adapting forms to specific contexts while reusing building blocks.

## Versioning and Lifecycle Implications

Every major artifact in the form service – fields, components, and forms – is **versioned and has lifecycle flags**. The schema uses a pattern of `<artifact>_business_key` plus a numeric version to identify different versions of the "same" conceptual artifact within a tenant <sup>23</sup> <sup>7</sup>. For example, a component with business key "customer\_address" might have version 1, 2, etc., each as separate records in the `component` table. Uniqueness constraints ensure no duplicate (`business_key`, `version`) within a tenant <sup>24</sup>. This indicates the system can retain old versions alongside new ones – crucial for a scenario like a published form that end-users are filling out (version 1 remains active for existing links or submissions) while a newer version 2 is available for new use cases.

To manage lifecycle, each artifact also has `is_published` and `is_archived` boolean flags (with corresponding timestamps) that control its state <sup>25</sup>. **Publishing** an artifact makes it available for use – e.g. a published field or component can appear in builder palettes for new forms <sup>25</sup>. **Archiving** marks it as retired from active use, though kept for historical purposes (the schema notes archived items are retained for back-compatibility of older form versions or submissions) <sup>26</sup>. The schema enforces consistency, for

example requiring that if `is_published` is true then `published_at` is not null, and similarly for archived <sup>27</sup>, which prevents logical errors. New entries default to published (true) and not archived (false) in the case of components <sup>25</sup>, implying that by default new artifacts are immediately usable unless explicitly hidden.

One subtle point in the schema is how versioning interacts with the stable runtime keys. For fields, `field_key` must be unique per tenant (it's the stable identifier used in forms) <sup>2</sup>. This suggests that if a field is versioned, the new version might have to either reuse the same `field_key` (which the unique constraint would normally forbid if the old version record still exists active) or the system updates the old version's key or usage status. The schema doesn't explicitly describe the process, but given the presence of a business key separate from `field_key`, it's likely the intent is that the stable `field_key` remains the same across versions, and only one version of a given field is meant to be "active" at a time (older versions might be archived to free the key). Similarly, components have a unique `component_key` per tenant <sup>28</sup>, separate from the versioned business key. In practice, this means a component concept (say "Address Block") might keep the same `component_key` for integration references, while its `component_business_key` + version distinguishes versions under the hood. The schema's constraints suggest that simultaneous active versions with the same key are not allowed, enforcing an **update-in-place or retire-old-then-add-new** version strategy for practical use. This ensures integrations or form submissions referencing a stable key aren't confused by duplicates.

The lifecycle flags combined with versioning provide a **form of governance over change**. For example, one could have a "draft" newer version of a form (`is_published = false`) that is being edited while an older version is still live. Once ready, the new version can be published and perhaps the old one archived. Old versions remain in the database, so submissions tied to them can still be understood (since the form definition record is retained as history) <sup>7</sup> <sup>29</sup>. This is a significant capability: it implies the platform is designed to avoid breaking changes – you don't delete or overwrite definitions that live submissions depend on; you version them.

## Tenant vs. Provider vs. Marketplace Separation

The schema is explicitly multi-tenant: every primary table includes a `tenant_id` to scope data, and **tenant-safe foreign keys** are used throughout to enforce that references don't cross tenant boundaries <sup>30</sup> <sup>31</sup>. This means each tenant (organization using the service) has their own isolated set of forms, fields, components, etc., which is typical for a SaaS. However, the design also anticipates **two other scopes of artifact origin**: provider and marketplace. Each artifact (field, component, form) has metadata columns (`source_type`, `source_package_key`, `source_artifact_key`, etc.) that indicate where it came from <sup>32</sup> <sup>33</sup>. The `source_type` is an enumeration of values like *TENANT*, *PROVIDER*, *MARKETPLACE*, *SYSTEM* <sup>32</sup>. This provides a clear separation between things the tenant created themselves versus things that were installed or provided.

- A **TENANT** artifact means it was created within that tenant's environment (their own custom field or form).
- **PROVIDER** likely refers to artifacts provided by the platform vendor or a specific package of defaults (for example, the system might ship with some standard fields or form templates that appear in each tenant).
- **MARKETPLACE** indicates the artifact was imported from a marketplace – presumably a third-party or community-contributed package of form definitions or components.

- **SYSTEM** could be reserved for internal usage artifacts.

Crucially, even when an artifact originates from a marketplace or provider, the schema stores it in the tenant's tables (with that tenant's ID) – it's a **local copy with provenance metadata**. The `source_package_key` and `source_artifact_key` act as identifiers to trace which package and item this artifact came from <sup>34</sup> <sup>35</sup>. This approach ensures *separation with control*: a marketplace component imported into Tenant A lives in Tenant A's data space (no runtime dependency on a global object), preventing cross-tenant data mixing and allowing Tenant A to modify it independently. At the same time, the original keys and a `source_checksum` are stored, so the system knows if Tenant A's copy has diverged from the source or if an update is available <sup>36</sup> <sup>37</sup>.

Additionally, the use of stable business keys across tenants suggests that if two tenants install the same marketplace component, they would share a common business key (perhaps the package's global identifier) but have distinct records. The schema comment explicitly notes that a component's business key and version can be used across tenants via marketplace import <sup>5</sup> – implying a strategy to align artifacts by keys even in different tenant databases. This would facilitate things like marketplace updates (knowing that tenant X's "Address Block" v1 corresponds to marketplace "Address Block" v1, for example).

In summary, the schema supports a **clear separation of tenant-owned vs externally-sourced content**, with all the necessary keys to manage distribution. There is no mingling of data at the row level, and all foreign keys are compound (`tenant_id`, `id`) pairs to avoid any reference escaping its tenant <sup>38</sup> <sup>39</sup>. This structure is well-aligned with multi-tenant security and also sets the stage for managing a marketplace of shareable form components safely.

## Governance and Control Surface

The schema implies several governance mechanisms for controlling how forms and their building blocks are used. First, the presence of **audit columns** (`created_by`, `updated_by`, timestamps on every table) indicates that tracking who created or modified an artifact is considered important <sup>40</sup> <sup>41</sup>. This is typical for enterprise governance – one can audit changes to forms or fields.

The **lifecycle flags (is\_published, is\_archived)** provide control over what is currently active. For example, an administrator can mark a field definition as not published to hide it from use without deleting it, or archive a form so it no longer appears for end-users. The schema's enforcement of consistency for these flags <sup>27</sup> <sup>42</sup> means the system will reliably reflect the intended state (there won't be "ghost" published items with missing timestamps, etc.). This gives a measure of control to roll out and roll back availability of form elements.

**Tenant isolation** is also a governance aspect – it ensures one tenant cannot accidentally reference another's form or field. The use of composite unique constraints like `ux_field_def_id_tenant` and matching composite foreign keys (`tenant_id`, `id`) in all relationships means the database will reject any cross-tenant linking <sup>38</sup> <sup>43</sup>. This acts as a safety net beyond application logic, preventing data leakage or mis-assignment.

The schema introduces a concept of **categories** (`form_catalog_category`) which are tenant-defined groupings for fields and components <sup>44</sup>. Categories have `is_active` flags <sup>45</sup>, meaning an admin can enable/disable entire groups of elements (perhaps to govern a palette of building blocks available to form

builders). They also have stable keys, suggesting they can align with marketplace categories or be used in import/export mapping <sup>46</sup>. The existence of categories and the ability to activate/deactivate them provides another layer of organizational control, helping administrators govern how a potentially large library of fields/components is presented and used.

Finally, the **panel\_actions** JSON on panels (both form panels and component panels) hints at a controlled surface for interactive behavior <sup>47</sup> <sup>48</sup>. While more of a functional feature than administrative governance, panel actions allow defining rules (show/hide logic, interdependencies, etc.) in a declarative way. This can be seen as governing the **user experience** – it gives form designers a way to enforce certain behaviors without coding. For example, a panel could be set to remain hidden unless a checkbox field is true, encoded in the JSON rules. This is a form of *governance of form logic* at the schema level.

In summary, the schema provides multiple control surfaces: **access control by tenant, state control via publish/archive, organizational control via categories, and auditability via created/updated metadata**. All of these point to a system designed to be managed in a disciplined way, suitable for enterprise use where oversight of form configurations and their evolution is needed. The schema itself doesn't include user roles or permissions, but it lays the groundwork (e.g. every table has `created_by` fields, implying integration with whatever identity system manages the platform's users). All these controls contribute to a governance model necessary for a form platform that may be extended by many parties (internal developers, third-party contributors, tenant admins) while maintaining order and security.

## Comparison to Leading CRM Platforms

When comparing this form service schema's capabilities to leading CRM platforms (like **HubSpot, Zendesk, Salesforce**), several differences in flexibility, design approach, and risk areas emerge:

- **HubSpot:** HubSpot's form builder is relatively straightforward and tightly tied to HubSpot's CRM data model. Each form field in HubSpot must map to a defined CRM property (contact or other object property) <sup>49</sup>. This means the universe of fields is global to the portal (tenant): you create a custom property once and then any form can use it. The Dyno form schema is similar in that it has global field definitions per tenant, but it goes further by allowing fields that aren't strictly CRM properties – they are form-specific data capture fields if desired (HubSpot lacks truly form-only fields; every field populates some property). HubSpot forms also support basic conditional logic and multi-step layouts, but **do not offer reusable components or versioned form definitions**. For example, you cannot define a sub-form in HubSpot and drop it into multiple forms; you would manually recreate those fields in each form. By contrast, the schema here enables complex reuse (via `component`) and even has a marketplace concept for distributing those components. HubSpot also doesn't version forms – editing a form is immediate and overwrites the old design (no concept of `form_v1` vs `form_v2` as separate records). This form service provides more rigorous change management with version histories, which is a strength in maintaining large or critical forms.

In terms of flexibility, HubSpot forms are limited by their dependency on CRM objects – you typically use them for lead capture or support tickets where data goes into HubSpot's contacts/tickets. The form service schema is **object-agnostic**: a form submission results in entries in `form_submission_value` which could then be mapped anywhere, giving the platform (or the customer) freedom to decide how form data is used. This is a level of flexibility more akin to dedicated form-building tools than a CRM's built-in forms. The trade-

off is complexity; HubSpot's simpler model is easier for average users and deeply integrated into their contact database, whereas this form service requires careful design of field definitions and mappings.

- **Zendesk:** Zendesk supports *ticket forms* in its support module. Administrators can define multiple ticket forms, but these are essentially different groupings of the same pool of global ticket fields. The model is: you create custom fields (global to the Zendesk account), then decide which forms those fields appear on <sup>50</sup>. This is conceptually similar to the field catalog in the Dyno schema – one set of field definitions, many forms – but Zendesk's approach is far more limited. It doesn't have the idea of nested components or sub-forms; a ticket form is one level list of fields (with some conditional display rules possible, e.g. show field B only if A has value X). The Dyno form service allows arbitrarily nested grouping and complex logic via panel actions, which goes well beyond Zendesk's capabilities. Additionally, **versioning** in Zendesk is essentially nonexistent: if you change a ticket form (add/remove fields), the change is immediately live; you would have to clone a form manually to version it. The Dyno schema's versioning would allow, for instance, keeping an old ticket form definition active for existing tickets while rolling out a new version for new tickets – something not readily achieved in Zendesk's standard features.

Another difference is **marketplace readiness**. Zendesk does not have a concept of sharing ticket forms or custom field sets across accounts via a marketplace. Each Zendesk instance's admins configure their fields/forms from scratch or by manually copying ideas. The Dyno form service, however, is explicitly built to import/export artifacts, suggesting a potential ecosystem of pre-built form templates or field/component libraries that users can leverage. This could be a competitive advantage in speed of setup or industry-specific solutions. The risk on Dyno's side might be managing these imports safely (which the schema addresses with provenance tracking and checksums). Zendesk avoids that complexity by simply not offering cross-tenant form sharing.

- **Salesforce:** Salesforce's core platform historically did not have a standalone form builder for external use; instead, forms are represented as page layouts or Lightning pages for internal users, or you use Web-to-Lead/Web-to-Case for simple external forms. Recently, Salesforce introduced **Dynamic Forms** (for Lightning record pages) and acquired capabilities like **OmniStudio** for advanced guided forms. Comparing to **Dynamic Forms**: Dynamic Forms let an admin design a record page by placing fields and sections with conditional visibility. This is somewhat analogous to designing a form with panels and fields. However, Dynamic Forms are bound to a Salesforce object (e.g., a Case or Lead) – you cannot use them to capture arbitrary data outside the predefined schema. If you need a form that doesn't map 1:1 to a single Salesforce object (Salesforce calls this "object-agnostic" or involving multiple objects), Dynamic Forms are *not available* <sup>51</sup> <sup>52</sup>. Salesforce would direct you to use a tool like a Flow or OmniStudio for those scenarios. In contrast, the Dyno form service is **inherently object-agnostic** – it captures data in a generic way (JSON values per field) and only later might you map those to actual CRM records. This is a design more flexible than standard Salesforce forms, closer to Salesforce's **OmniStudio** offering. In fact, OmniStudio (now part of Salesforce) uses a JSON metadata approach where a form (OmniScript) is defined independent of the database and can then be mapped to objects or APIs after submission <sup>53</sup>. The Dyno form schema aligns with that philosophy: forms defined in JSON and metadata, storing results in a generic table, and presumably the application layer or an integration will decide what to do with the results (create records, send to external service, etc.). In essence, Dyno's form service could be seen as offering Salesforce OmniStudio-like power within a CRM platform.

**Versioning and packaging** are areas where this form service appears stronger than Salesforce's native tools. Salesforce does have packaging for distributing customizations (managed packages on the AppExchange), but it's a heavier-weight process (involving deploying metadata to orgs). The Dyno marketplace concept would allow distribution of form components in a more "app-store" style, possibly with one-click install into a tenant. This lowers the barrier for sharing and could outpace Salesforce in flexibility if executed well. On the flip side, Salesforce's platform is very mature in governance – changes are tracked, and admin oversight is high. Dyno's approach will need equally strong UI/process around these features to ensure non-developers can manage the complexity (the schema provides the capability, but the product must expose it simply).

- **Others:** Other CRM or support platforms (e.g., Microsoft Dynamics, ServiceNow, Freshdesk) each have their own form or field customization approaches. Generally, most follow the pattern of *global field definitions and per-form selection*, with varying support for conditional logic. Very few offer the kind of **nested reusable component** paradigm we see in this schema. For instance, ServiceNow allows designing forms for its records but doesn't have marketplace form components in the way Dyno does (ServiceNow has an app store for entire applications, not for form snippets). The **flexibility of Dyno's form service is closer to application-building platforms (like Outsystems or Mendix, or form-specific platforms like Form.io)** than to traditional CRMs. This gives it a notable differentiation: it's trying to provide the extensibility of a platform while focusing on the domain of forms within a CRM context.

**Relative strengths:** The Dyno form service, as defined by the schema, is extremely flexible and extensible compared to typical CRM form builders. It supports reuse at scale (fields and components that can be shared), structured composition (hierarchies of panels), and safe evolution of forms (versioning and lifecycle flags). It enables a marketplace ecosystem, which leading CRMs do not natively have for forms (Salesforce comes closest via AppExchange, but again not at the granularity of individual form components). For customers who need highly customized processes or who want to rapidly adopt best-practice forms from a marketplace, this design is very appealing.

**Relative weaknesses/risks:** With great flexibility comes complexity. A potential risk is that this form service might be **over-engineered for simpler use cases**. Many HubSpot or Zendesk customers just need a basic form to collect a few fields – the overhead of managing field definitions, business keys, versions, etc., could be daunting if not abstracted by the UI. There's also performance and complexity considerations: storing each form submission as multiple rows (one per field) and heavy use of JSONB means reporting and querying might require more effort (the schema does index things like `field_path` and `field_def` for analytics<sup>54</sup> <sup>55</sup>, which mitigates this). A monolithic CRM form builder that writes directly to object fields might be simpler to integrate with the reporting and automation engines of that CRM. Dyno will need to ensure that the powerful features (like nested overrides, JSON merges, version branching) are harnessed in a user-friendly way, otherwise the learning curve could be a barrier.

Architecturally, Dyno's form service seems to be a **microservice or module separate from the core data model** (evidenced by its generic nature and tenant isolation). CRM platforms like Salesforce or HubSpot embed forms into their core (tightly coupling form fields to object fields). The decoupled approach gives Dyno more agility – forms can exist independently of CRM objects – but it also means **integration effort** is needed to connect form submissions to CRM records or workflows. Salesforce can assume a form's fields directly update a Lead record, for example, whereas Dyno's form submission might require a script or flow to take the submission data and create a Lead. This is a trade-off between *hard-wired simplicity* (Salesforce/HubSpot) and *general-purpose flexibility* (Dyno).

## Marketplace-Readiness of the Schema

The provided schema was explicitly augmented for “marketplace-grade” capabilities <sup>56</sup>, and it shows. Several design elements support safe reuse, distribution, and evolution of form artifacts in a marketplace context:

- **Provenance and Metadata:** Each major artifact table (fields, components, forms) has columns to record its source type and origin keys <sup>32</sup> <sup>35</sup>. This means that when an item is installed from a marketplace or provided by the platform, the system knows exactly where it came from (which package and what version). This is critical for a robust marketplace – it enables updates (the system can check if a new version of that package is available) and attribution (knowing which vendor or source provided it). The schema even stores checksums of the source artifact versus the current artifact <sup>36</sup> <sup>37</sup>, which is an advanced feature to detect local modifications or drift. For example, if a tenant imports a component and then edits it locally, the `current_checksum` will differ from the original `source_checksum` <sup>36</sup>. This way, the platform could warn the admin “you have modified this component; updating from marketplace might overwrite your changes” – a safe-guard for marketplace content.
- **Stable Identifiers for Alignment:** The use of human-stable keys (e.g., `field_def_business_key`, `component_business_key`, `category_key`) ensures that artifacts can be matched across tenants or against a global catalog <sup>46</sup> <sup>5</sup>. A marketplace package would likely define a fixed key for each artifact it contains. When imported, a tenant’s artifact retains that key (as the `business_key`), allowing the marketplace to later identify “Artifact X in tenant is version 1, the marketplace has version 2 available.” The unique constraint on (`tenant, business_key, version`) prevents key collisions within one tenant, and the presence of `source_package_key` ensures that even if two different packages coincidentally use the same business key for different artifacts, the system can distinguish them by package scope.
- **Imprinting and Safe Reuse:** A standout feature for safe reuse is the *imprinting mechanism* for field definitions used in components/forms. When a `field_def` is placed into a component (`component_panel_field`), the schema stores a **snapshot of the field definition JSON** at that time (`field_config` JSONB) along with a hash of the source definition <sup>57</sup> <sup>58</sup>. This means the component has its own self-contained definition of the field as it was when added – insulating it from future changes to the global field definition. If the original `field_def` is updated later (say, new options added to a dropdown), the component’s snapshot remains as it was, unless an admin decides to re-import or update it. The schema can detect that the original changed because it keeps `source_field_def_hash` from the time of imprint <sup>58</sup>. This design shows a deep consideration for **evolution**: it’s akin to “copy by value” with the option to refresh. It prevents unintended side effects where updating a field in the library could break a component or form that relied on the old behavior. Safe reuse is achieved because each usage has a frozen copy of the needed specs, and **controlled updates** can be done when appropriate.
- **Isolation with the ability to Merge Changes:** The schema’s JSON merge functions <sup>21</sup> <sup>59</sup> play a role in how marketplace updates might be applied. Suppose a new version of a component is imported; if a tenant has local overrides (`nested_overrides`) for a particular instance, or even direct edits to the component), merging in the new default while preserving overrides is important.

The deterministic merge ensures that applying patches (like marketplace-driven updates) yields consistent outcomes. This is an implied capability: the tools are there to programmatically handle merging JSON configurations, which would be essential if, for example, a marketplace update wants to add a new field to a component without clobbering tenant-specific UI tweaks on existing instances.

- **Missing or Future Capabilities:** While the schema checks most boxes, a few marketplace-related features are not explicitly described, likely handled at the application level or future work:
  - There's no explicit notion of *dependency management* or *package version compatibility* beyond the artifact version itself. For instance, if a component from marketplace version 2 requires a new version of a related field also to be updated, the schema doesn't enforce that linkage (other than recording each artifact's own version). We assume the marketplace packaging process would handle such dependencies externally.
  - The schema doesn't have a dedicated table for marketplace packages or a manifest of installed packages. It treats each artifact individually. This is a simpler approach but means to list "what marketplace packages do I have installed?" the app might have to query distinct package\_key values across artifacts. A more consolidated representation (like an installed\_package table) is not present. This isn't a show-stopper, but it's an area for the application logic to cover.
  - **Security and Trust:** The schema itself cannot ensure that a marketplace artifact is safe or allowed for a tenant – that would be an external concern. But one might expect additional fields in the future like a signature or origin verification; currently, checksum covers integrity, not authenticity. However, those details are beyond schema and more about process/policy.
  - Another implied capability is **roll-back or parallel version installation:** The schema could support keeping an old and new version of a component concurrently (since version is part of the key). But whether the application will expose that (or force an upgrade path) is to be seen. The groundwork is there for having multiple versions, so presumably a robust marketplace could allow testing a new version of an artifact while still using the old one.

Overall, the schema appears **marketplace-ready**. It was clearly hardened for that use: we see thought given to how to import (keys, source tracking), how to update (checksums, versions), and how to allow local customization safely (imprints and overrides). This design minimizes risk in reuse: tenants get copies they fully control, but with references back to the original for updates. A missing piece in many simpler form builders – the ability to import a template and then diverge from it without losing the link – is directly addressed here.

## Strategic Positioning

Given these capabilities, the Dyno Form Service schema positions the product as an **extensible, platform-like form builder** within the CRM or ticketing domain. It is not a basic feature; it's a strategic differentiator aimed at power users, enterprises, and an ecosystem of partners. Key strategic implications include:

### Alignment with a Market-Leading, Extensible Form Platform

The schema is aligned with what we'd expect in a market-leading form platform, perhaps even overshooting what current CRM vendors offer. It provides the kind of **metadata-driven, modular architecture** that one finds in highly extensible platforms (for example, Salesforce's metadata model or ServiceNow's application

engine) rather than a typical hard-coded form builder. The inclusion of versioning, marketplace distribution, and multi-tenant considerations shows an intent to enable a **community and longevity** around form components. This is similar to how Salesforce enabled AppExchange or how Atlassian enables plugin marketplaces – by having stable extension points and version management. In the context of forms, this is forward-thinking and could make the product attractive to customers who want to avoid being locked into one vendor's limited form tool. Instead, they get a platform where they can build, share, and continuously improve forms and form elements.

The schema's alignment with **object-agnostic data capture** also means it can serve as a front-end for many scenarios. It's not just aligned with CRM use-cases (like lead forms or ticket forms); it can potentially address any workflow that needs a form. This breadth means the platform can position itself not only against CRM forms, but also against dedicated form/survey tools (Typeform, Google Forms, etc.) when integrated properly. Few of those tools offer the level of reusability and integration that Dyno's form service can (e.g., you can't easily version a Typeform or inject a pre-built "address block" module into it). So strategically, Dyno's form service could pitch itself as *the last form builder you need*, one that combines the ease-of-use of form apps with the enterprise integration of a CRM.

## Ideal Customer Types and Use Cases

The capabilities best support **enterprise and mid-market customers** who have complex or dynamic data collection needs. For example:

- **Organizations with multiple brands or products:** They can create a core library of fields/components and then compose tailored forms for each brand, while maintaining consistency. The tenant structure covers isolation at the org level, and categories could organize by product lines internally.
- **Industries with regulated or evolving processes:** such as insurance or healthcare where forms (applications, claims, intake forms) frequently change. The versioning ensures they can update forms while keeping audit trails of old versions (important for compliance).
- **Customers that require multiple forms serving different purposes:** A single company might need contact forms, support request forms, onboarding forms, feedback surveys – all of which could benefit from common elements (like contact info fields, address, etc.). This form service would let them avoid redefining fields each time and easily roll out changes across forms (update the component or field once, propagate via marketplace update or consistent keys).
- **Managed service providers or SaaS resellers:** If Dyno were used by a provider that serves multiple client orgs, the marketplace model allows a "provider" to develop form templates centrally and distribute them to clients (source\_type "PROVIDER"). For example, a consultancy could build a set of best-practice forms and deliver to all its client tenants on the platform. This multi-tenant replication use-case is directly supported by the schema design <sup>60</sup>.

Conversely, small businesses with very simple needs might find this system overkill if exposed directly – but Dyno could abstract complexity for them by providing out-of-the-box templates (leveraging the same marketplace system). Those smaller customers would benefit from the results (high-quality forms) without necessarily using all features. However, the real sweet spot is customers who treat forms as critical assets – those who would appreciate lifecycle management, reuse, and integration. For instance, a government agency using the platform could have dozens of forms that need strict version control and sharing of standard components (like a privacy statement section), which is exactly the scenario this schema handles well.

## Trade-offs vs. Monolithic CRM Form Builders

Compared to monolithic form builders embedded in CRMs, Dyno's approach trades simplicity for extensibility. Monolithic builders (e.g., a basic "web-to-lead" form generator) are quick to deploy but very inflexible – you get a fixed set of fields and minimal logic. Dyno's form service sacrifices that quick simplicity; it requires design upfront – defining field defs, maybe setting up components – but yields a much more flexible outcome. The **initial setup cost** is higher, but the **long-term payoff** is easier maintenance and scaling. If a monolithic form builder is a single-use tool, Dyno's is a platform capability that grows with the organization.

One trade-off is that Dyno's form submissions are stored in a generic way (JSON values per field), whereas monolithic forms often store data directly in first-class CRM records. This means out-of-the-box reporting or automation might require an extra step for Dyno's forms (to translate submission data into, say, a Contact record or a custom object). But this trade-off is intentional to decouple the form from the data model – making the forms more reusable and able to handle use cases where you might not even create a CRM record (perhaps the form triggers an external process instead).

Another consideration is **performance and complexity** in the UI: a drag-and-drop form builder that exposes all these features (fields, components, overrides, logic) will be complex. Market-leading CRMs have spent years refining simpler form editors for admins. Dyno will need to ensure the power features can be progressively revealed – an admin could do basic things easily, and only dive into components/overrides if they need to. If done well, this positions Dyno as a solution that **starts easy but can go deep**, which is compelling against competitors: you don't outgrow it when you hit complexity, whereas you might outgrow HubSpot or Zendesk forms and need a custom solution.

Strategically, the schema shows Dyno is aiming to be not just a form builder, but an **extensible platform within the platform**. It's the kind of design one would create if the goal is to allow third-party developers and power-users to extend the CRM's capabilities in unforeseen ways (much like Salesforce's early strategy with a strong metadata core that enabled AppExchange apps). If executed properly, this could foster a community around Dyno's product, driving adoption through available templates and components. The clear separation of tenant and provider content means it's ready to handle a multi-tenant cloud service at scale, another sign of market-leading ambition (some smaller CRM players only support single-tenant or lack true isolation, which becomes an issue in SaaS scaling; Dyno is avoiding that from the get-go).

In conclusion, the Postgres schema reveals a form service with **strengths in flexibility, reusability, and governance that surpass many traditional CRM form tools**. Its marketplace-oriented enhancements indicate a vision for a broader ecosystem. The trade-offs lie in added complexity and the need for strong management UX, but these are solvable with good product design. Dyno's form service appears well-positioned to serve complex enterprise needs and differentiate itself as a platform for forms – turning what is often a simple feature into a strategic asset. If monolithic CRM form builders are one-size-fits-all, Dyno's is a tailor-ready wardrobe, with all the tailoring tools provided.

---

1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30  
31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47 48 54 55 56 57 58 59 60

**form\_schema\_updated.sql**

file:///file-3p1wmBVuuBXcuCkP3mQT5

**49 Create and edit forms**

<https://knowledge.hubspot.com/forms/create-and-edit-forms>

**50 Adding custom ticket fields to your tickets and forms – Zendesk help**

<https://support.zendesk.com/hc/en-us/articles/4408883152794-Adding-custom-ticket-fields-to-your-tickets-and-forms>

**51 52 53 Building Forms | Salesforce Architects**

<https://architect.salesforce.com/decision-guides/build-forms>